

```
In [25]: 1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split, cross_val_score
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.pipeline import Pipeline
6 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier, ExtraTreesClassifier
7 from sklearn.svm import SVC
8 from sklearn.linear_model import LogisticRegression, RidgeClassifier
9 from sklearn.neighbors import KNeighborsClassifier
10 from sklearn.tree import DecisionTreeClassifier
11 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, AdaBoostRegressor, ExtraTreesRegressor
12 from sklearn.linear_model import LinearRegression, Ridge
13 from sklearn.metrics import accuracy_score, classification_report, mean_absolute_error, mean_squared_error, r2_score

In [26]: 1 # Load dataset (Wine Quality dataset from UCI Repository)
2 df = pd.read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv", sep=
3
4

In [27]: 1 #define features and target
2 X = df.drop(columns=['quality'])
3 y = df['quality']

In [28]: 1 # Convert target into binary classification (Good vs. Bad wine)
2 y_class = y.apply(lambda x: 1 if x >= 6 else 0)
3
4 # Split the dataset for classification
5 X_train_c, X_test_c, y_train_c, y_test_c = train_test_split(X, y_class, test_size=0.2, random_state=42)
6
7

In [29]: 1 #classification models
2 classification_models = {
3     'Logistic Regression' : LogisticRegression(),
4     'Ridge Regression' : RidgeClassifier(),
5     'Random Forest' : RandomForestClassifier(n_estimators=50, random_state=41),
6     'Gradient Boosting' : GradientBoostingClassifier(n_estimators=100, random_state=42),
7     'AdaBoost' : AdaBoostClassifier(n_estimators=50, random_state=41),
8     'Extra Trees' : ExtraTreesClassifier(n_estimators=50, random_state=41),
9     'SVM' : SVC(kernel='rbf', probability=True),
10    'K-Nearest Neighbours' : KNeighborsClassifier(n_neighbors=5),
11    'Decision Tree' : DecisionTreeClassifier(random_state=41)
12
13 }
```

```

In [30]: 1 best_classification_model = None
2 best_classification_accuracy = 0
3
4 # Iterate over classification models and evaluate
5 print("\nClassification Models Evaluation")
6 for name, model in classification_models.items():
7     pipeline = Pipeline([
8         ('scaler', StandardScaler()),
9         ('classifier', model)
10    ])
11
12    # Train the model
13    pipeline.fit(X_train_c, y_train_c)
14
15    # Predict
16    y_pred_c = pipeline.predict(X_test_c)
17
18    # Evaluation metrics
19    accuracy = accuracy_score(y_test_c, y_pred_c)
20    report = classification_report(y_test_c, y_pred_c)
21
22    print(f"\n{name} Classification Performance:")
23    print(f"Accuracy: {accuracy:.4f}")
24    print("Classification Report:")
25    print(report)
26
27    if accuracy > best_classification_accuracy:
28        best_classification_accuracy = accuracy
29        best_classification_model = name
30
31

```

```

Classification Report:
              precision    recall  f1-score   support

     0       0.80      0.80      0.80       141
     1       0.84      0.84      0.84       179

 accuracy          0.82          0.82          0.82       320
 macro avg         0.82          0.82          0.82       320
 weighted avg      0.82          0.82          0.82       320

```

SVM Classification Performance:

Accuracy: 0.7719

Classification Report:

```

              precision    recall  f1-score   support

     0       0.73      0.77      0.75       141
     1       0.81      0.77      0.79       179

 accuracy          0.77          0.77          0.77       320

```

```

In [31]: 1 #split for regression
2
3 X_train_r,X_test_r,y_train_r,y_test_r = train_test_split(X, y, test_size=0.25, random_state=41)
4
5 #reg models
6 regression_models = {
7     'Linear Regression' : LinearRegression(),
8     'Ridge Regression' : Ridge(),
9     'Random Forest Regressor' : RandomForestRegressor(n_estimators=100, random_state=41),
10    'Gradient Boosting Regressor' : GradientBoostingRegressor(n_estimators=50, random_state=41),
11    'AdaBoost Regressor' : AdaBoostRegressor(n_estimators=50, random_state=41),
12    'Extra Trees Regressor': ExtraTreesRegressor(n_estimators=50, random_state=41)
13
14 }

```

```

In [37]: 1 best_regression_model = None
2 best_r2_score = -float("inf")
3
4 # Iterate over regression models and evaluate
5 print("\nRegression Models Evaluation")
6 for name, model in regression_models.items():
7     pipeline = Pipeline([
8         ('scaler', StandardScaler()),
9         ('regressor', model)
10    ])
11
12    # Train the model
13    pipeline.fit(X_train_r, y_train_r)
14
15    # Predict
16    y_pred_r = pipeline.predict(X_test_r)
17
18    # Evaluation metrics
19    mae = mean_absolute_error(y_test_r, y_pred_r)
20    mse = mean_squared_error(y_test_r, y_pred_r)
21    r2 = r2_score(y_test_r, y_pred_r)
22
23    print(f"\n{name} Regression Performance:")
24    print(f"Mean Absolute Error: {mae:.4f}")
25    print(f"Mean Squared Error: {mse:.4f}")
26    print(f"R^2 Score: {r2:.4f}")
27
28    if r2 > best_r2_score:
29        best_r2_score = r2
30        best_regression_model = name
31
32

```

Regression Models Evaluation

Linear Regression Regression Performance:

Mean Absolute Error: 0.5032
Mean Squared Error: 0.4199
R^2 Score: 0.3619

Ridge Regression Regression Performance:

Mean Absolute Error: 0.5033
Mean Squared Error: 0.4199
R^2 Score: 0.3619

Random Forest Regressor Regression Performance:

Mean Absolute Error: 0.3964
Mean Squared Error: 0.3032
R^2 Score: 0.5392

Gradient Boosting Regressor Regression Performance:

Mean Absolute Error: 0.4623
Mean Squared Error: 0.3519
R^2 Score: 0.4652

AdaBoost Regressor Regression Performance:

Mean Absolute Error: 0.4758
Mean Squared Error: 0.3838
R^2 Score: 0.4166

Extra Trees Regressor Regression Performance:

Mean Absolute Error: 0.3511
Mean Squared Error: 0.2763
R^2 Score: 0.5800

```

In [38]: 1 # Conclusion
2 print("\nConclusion:")
3 print(f"The best classification model is {best_classification_model} with an accuracy of {best_classification_accu")
4 print(f"The best regression model is {best_regression_model} with an R^2 score of {best_r2_score:.4f}.")

```

Conclusion:

The best classification model is Extra Trees with an accuracy of 0.8219.
The best regression model is Extra Trees Regressor with an R^2 score of 0.5800.

```

In [ ]: 1

```